

SWORAJ SAHOO

Analyst AI/LLM | Generative AI Developer

Phone: 7205235887 | Email: sworajsahoo05@gmail.com | [linkedin.com/in/sworaj-sahoo-6789a0273](https://www.linkedin.com/in/sworaj-sahoo-6789a0273)

PROFESSIONAL SUMMARY

Results-driven AI professional with hands-on experience in Generative AI development, LLM engineering, and AI training data operations. Currently working as Analyst AI/LLM at Innodata, building RAG pipelines, LLM-powered applications, and prompt engineering solutions. Prior experience at GlobalLogic as an Associate Analyst honed deep expertise in data annotation, multilingual content moderation, and Trust & Safety policy enforcement for large-scale AI training datasets. Proficient in Python, LangChain, LlamaIndex, OpenAI API, Anthropic Claude API, and vector databases including Pinecone and ChromaDB.

TECHNICAL SKILLS

Core AI/ML: Retrieval-Augmented Generation (RAG), Prompt Engineering, Chain-of-Thought (CoT), Few-Shot Prompting, Text Embeddings, Vector Search, Fine-Tuning, LoRA, PEFT

Frameworks & Libraries: LangChain, LlamaIndex, Hugging Face Transformers, PEFT, pandas, numpy, asyncio, FastAPI, pydantic

LLM APIs & Models: OpenAI API (GPT-4, GPT-3.5), Anthropic Claude API, Google GenAI SDK, Llama 2, Mistral 7B, VGG19

Vector Databases: Pinecone, ChromaDB, Milvus

Programming Languages: Python,

Tools & Platforms: Git, GitHub, VS Code, Jupyter Notebooks, Streamlit, Gradio, Weights & Biases, Docker (basic)

Data & AI Operations: Data Annotation, Content Labeling, AI Training Datasets, Trust & Safety, Content Moderation, Inter-Annotator Agreement

WORK EXPERIENCE

Analyst AI/LLM | Innodata

January 2026 – Present | Noida, India

- Design and deploy Retrieval-Augmented Generation (RAG) pipelines connecting large language models (LLMs) to private document stores using LangChain and ChromaDB to ground responses in verified data and reduce hallucinations.
- Engineer advanced prompt templates using Chain-of-Thought (CoT) and Few-Shot prompting strategies, improving model output accuracy and consistency for complex reasoning tasks.
- Build end-to-end AI chatbot and question-answering pipelines using the OpenAI API and Anthropic Claude API, implementing streaming responses, async handling, and structured JSON output parsing.
- Implement text embedding workflows to convert documents into high-dimensional vectors stored in Pinecone, enabling efficient semantic search and retrieval for LLM applications.
- Experiment with parameter-efficient fine-tuning (LoRA/PEFT) on open-source models including Llama 2 and Mistral 7B via Hugging Face to adapt models for domain-specific tasks.
- Develop interactive demo frontends using Streamlit and Gradio to present AI application prototypes to internal stakeholders.
- Optimize token usage and context window management to reduce API costs through intelligent document chunking and response caching strategies.
- Implement system-level guardrails and prompt injection defenses to maintain safe and policy-compliant model outputs in production environments.

Associate Analyst | GlobalLogic

November 2024 – December 2025 | Gurugram, India

- Performed data annotation and content-labelling tasks for large-scale AI training datasets, directly contributing to model quality and dataset accuracy enhancement.

- Annotated and categorized multilingual text, image, and video content to support AI model training pipelines and large-scale dataset quality improvement initiatives.
- Reviewed, identified, and flagged sensitive and harmful content including hate speech, sexual content, harassment, violence, and self-harm indicators in strict accordance with Trust & Safety policies.
- Applied and upheld complex annotation guidelines to ensure accuracy, consistency, and policy alignment across diverse content labelling tasks.
- Participated in quality-calibration sessions to maintain consistent judgment and improve inter-annotator agreement across the content moderation team.
- Ensured full compliance with Trust & Safety standards while consistently meeting high labelling accuracy and productivity targets.

ACADEMIC PROJECTS

Brain Tumour Detection | Machine Learning, VGG19, Transfer Learning

- Developed a brain tumour detection model leveraging VGG19 transfer learning architecture for medical image classification.
- Implemented complete data preprocessing, model training, and evaluation pipelines, achieving improved diagnostic accuracy on MRI image datasets.
- Applied image augmentation and hyperparameter tuning techniques to enhance model generalization and reduce overfitting.

Weather Forecast Application | Python, Django, REST APIs

- Built a full-stack weather forecasting application integrating data from multiple external weather APIs to provide real-time and accurate forecasts.
- Designed backend logic using Python and Django to store and manage weather data in a centralized relational database.
- Developed user-facing features for location-based weather lookup, multi-day forecasting, and historical weather trends.

EDUCATION

Master of Computer Applications (MCA)

K.R. Mangalam University, Gurugram | 2022 – 2024

Bachelor of Science (BSc)

Utkal University, Bhubaneswar | 2019 – 2022

CERTIFICATIONS AND TRAINING

- GitHub Foundations Certification — GitHub (March 2025 – March 2028)
- Python Full Stack Training — Aptiron Solutions